



Bolighed - powered by Python

A real life case study

About me



- Mathematician by education
- After some years in research I have since worked as a Python developer, primarily in data processing at:
 - Danish Geodata Agency, now called SDFE
 - Danish Meteorological Institute
 - And now, **Bolighed A/S**

About Bolighed



- [Bolighed](#) is a website aimed at house owners and hunters as well.
- Collects and presents information from a lot of public data sources:
 - BBR (basic information about buildings, dwellings and ownership).
 - Tinglysning (loans, entitlements etc.)
 - Energy marks
 - ...
- Also a lot of “closed” non-public sources:
 - Price estimate models (machine learning)
 - Sales data
 - ...

The stack at Bolighed

Advanced setup with a **lot** of components:

- Amazon EC2
- Docker
- Redis
- Elasticsearch
- Cloudflare
- Postgres / Postgis (databases)
- Nginx
- Tornado
- ... and a whole lot more ...

Where is Python used?



The frontend is using AngularJS and Python takes care of the rest:

- Data import
- Infrastructure
 - Deployment / configuration via **ansible**
- Backend api
 - Flask
 - Django
- Data analysis
 - Numpy, scipy, pandas, matplotlib, scikit-learn, SQL via SQLAlchemy.

A deeper look into some of the use cases

Backend:

- Flask with SQLAlchemy
- Super simple and very flexible setup:



```
from flask import Flask  
app = Flask(__name__)
```

```
@app.route('/')  
def hello_world():  
    return 'Hello, World!'
```

Backend

Why Python and not PHP, C, C# or Java (or Ruby)? Is performance OK??

- Much, much nicer and more maintainable than PHP!
- Very high level interface to various services / infrastructure
 - Elasticsearch, Redis, Postgres (SQLAlchemy), Datadog, Amazon EC2 / S3.
- Lot's of caching mechanisms and load balancing in place - very few actual database calls...

Backend

We also have some apis running in Django:



- More structured than Flask + SQLAlchemy
- Includes it's own ORM (Object-relational mapping) as a high level interface to the database.
- Lot's of extensions, e.g. The Django REST framework logo, with 'django' in small text above 'REST' in large outlined letters, and 'framework' in small text below.
- Used by many *huge* web applications out there:
 - **Instagram**
 - **Pinterest**
 - ...



Django's ORM

```
class CustomerType(models.Model):
    created = models.DateTimeField(auto_now_add=True)
    modified = models.DateTimeField(auto_now=True)
    name = models.CharField(max_length=255, unique=True)
    def __str__(self):
        return self.name

class PropertyData(models.Model):
    """
    Models any kind of property
    """
    bbr_property_data = models.ForeignKey('BBRPropertyData', null=True)
    address = models.ForeignKey('Address', null=True)
```

Django's ORM

```
(venv_bm) Simons-MacBook-Pro:business_manager simonkockendorff$ python manage.py shell
```

```
Python 3.6.0 (default, Dec 24 2016, 08:01:42)
```

```
Type "copyright", "credits" or "license" for more information.
```

```
...
```

```
In [1]: from business_manager.leads import models
```

```
In [2]: for obj in models.Address.objects.all().filter(street__startswith="Åsvej")[:2]:
```

```
....:     print(obj)
```

```
....:
```

```
Åsvejen 4 , 4330
```

```
Åsvejen 6 , 4330
```

- Specific database is 'abstracted away'
- No explicit SQL queries
- However, in some cases the high level ORM is too rigid and one must resolve to plain old SQL...

Data import

We use a lot of different python libraries and protocols for fetching data from various sources:

- **Boto / boto3** for talking to Amazon EC2 and S3
- **Requests** for REST-interfaces / scraping
- **Pysimplesoap / Requests** for SOAP (XML) interfaces (sigh....)

For example there is a great API for all danish addresses at <http://dawa.aws.dk/>

```
In [14]: import requests
```

```
In [15]: r = requests.get("http://dawa.aws.dk/adresser", params={"vejnavn": "Fasanvej", "postnr": 8210, "husnr": 15, "struktur": "mini"})
```

```
In [16]: r.json()
```

```
Out[16]:
```

```
[{'adgangsadresseid': '0a3f5096-212e-32b8-e044-0003ba298018',  
  'dør': None,  
  'etage': None,  
  'husnr': '15',  
  'id': '19910d90-1d47-41c9-e044-0003ba298018',  
  'kommunekode': '0751',  
  'postnr': '8210',  
  'postnrnavn': 'Aarhus V',  
  'status': 1,  
  'supplerendebynavn': None,  
  'vejkode': '2032',  
  'vejnavn': 'Fasanvej',  
  'x': 10.1787079932534,  
  'y': 56.1647588529531}]
```

Addresses, postal districts and various other data are imported from this endpoint on a regular basis.

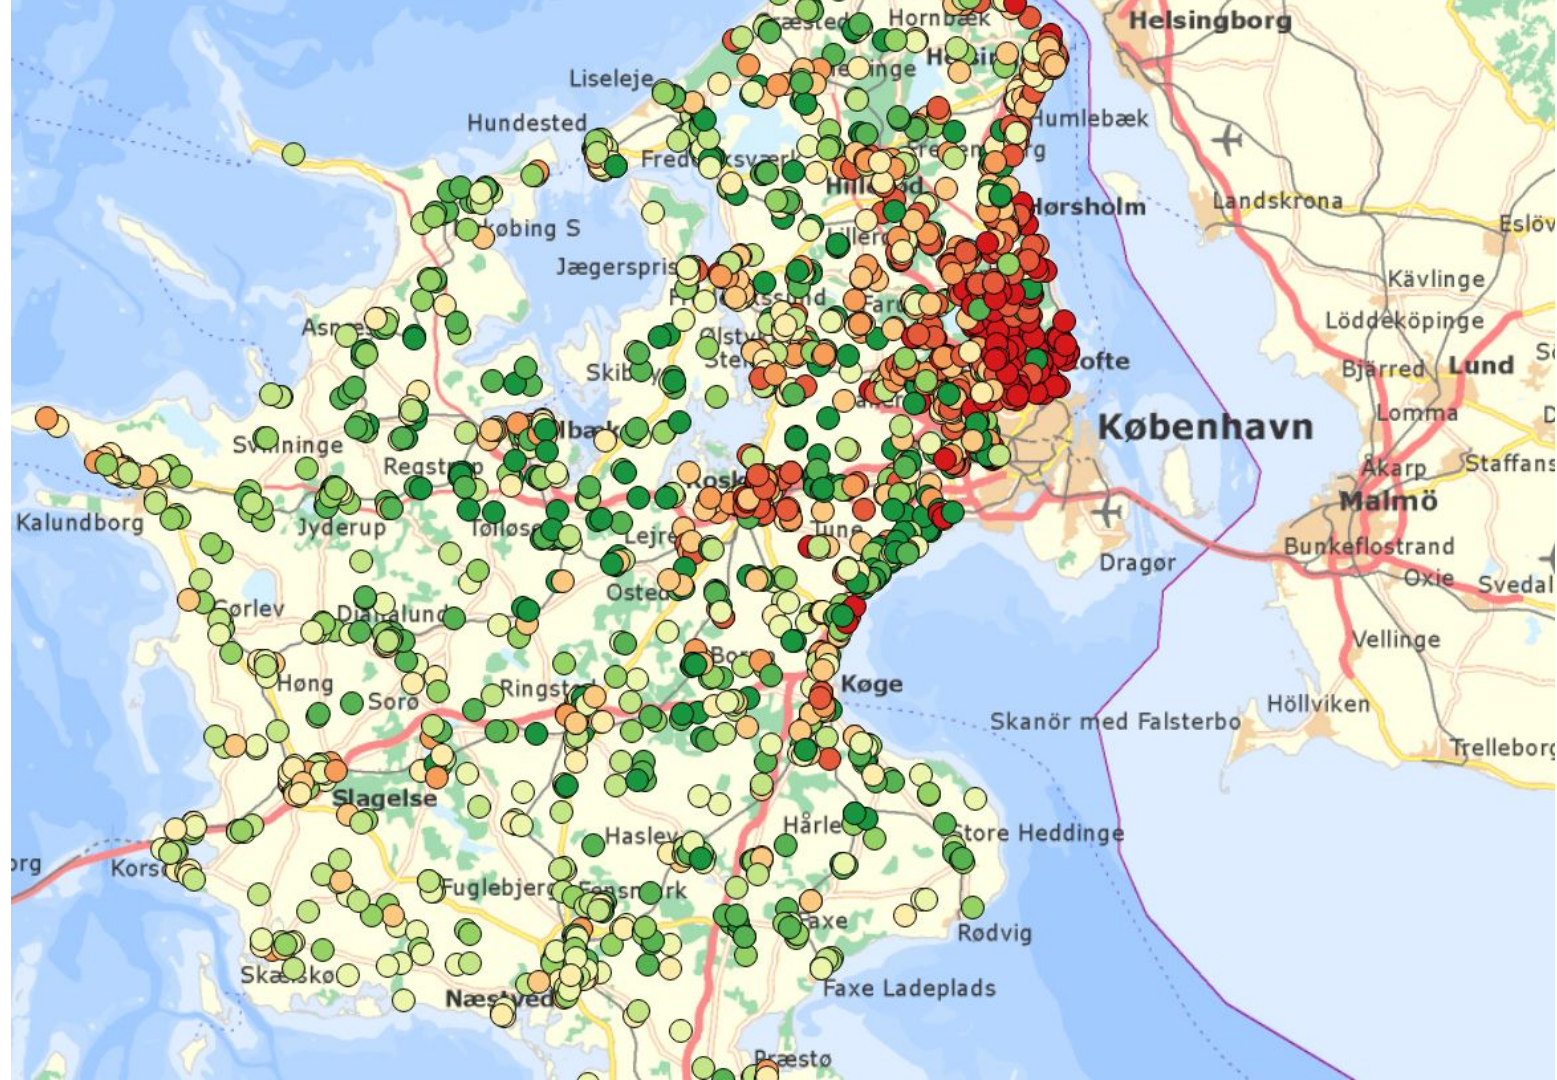
Data analysis

Case:


- Examine what the relation between house prices and travel time to Copenhagen is?

Plan:

- Fetch sales data + geographic location from database (Postgis) via SQLAlchemy.
- Use googlemaps Python API to query travel times to Copenhagen Central station for these locations.
- Do some analysis and plotting with numpy (linear regression, filtering) and matplotlib



kbh/kbh.csv



Something else that I've been working on...

- Mapping value increases for houses the next year:
 - <https://s3.bolighed.dk/static/stories/prisprognose/index.html#7/56.188/11.646>
- And something completely different - a fancy map:
 - <http://gittebach.dk/case/story.html>
- How does house prices depend on various parameters?
 - For example energy marks?
 - Create models using scikit-learn...
 - ...or tensorflow ... or...

Work in progress... analysis with statsmodels

OLS Regression Results

```
=====
Dep. Variable:          sqm_price    R-squared:                0.030
Model:                  OLS          Adj. R-squared:            0.029
Method:                 Least Squares  F-statistic:              38.69
Date:                  Tue, 28 Mar 2017  Prob (F-statistic):       7.19e-54
Time:                  14:46:46       Log-Likelihood:          -94526.
No. Observations:      8890          AIC:                    1.891e+05
Df Residuals:          8882          BIC:                    1.891e+05
Df Model:              7
Covariance Type:       nonrobust
=====
```

	coef	std err	t	P> t	[0.025	0.975]
Intercept	1.569e+04	649.095	24.176	0.000	1.44e+04	1.7e+04
energy_mark[T.C]	-1227.6587	561.354	-2.187	0.029	-2328.043	-127.274
energy_mark[T.D]	-1839.8345	549.096	-3.351	0.001	-2916.190	-763.479
energy_mark[T.E]	-2823.8346	569.063	-4.962	0.000	-3939.330	-1708.339
energy_mark[T.F]	-3690.3904	608.908	-6.061	0.000	-4883.991	-2496.789
energy_mark[T.G]	-7204.4471	631.368	-11.411	0.000	-8442.075	-5966.819
energy_mark[T.H]	-1.362e+04	1.01e+04	-1.355	0.176	-3.33e+04	6088.004
room_count	143.3920	74.239	1.931	0.053	-2.134	288.918

```
=====
Omnibus:                4939.886    Durbin-Watson:            0.701
Prob(Omnibus):           0.000      Jarque-Bera (JB):         64405.615
Skew:                    2.398      Prob(JB):                 0.00
Kurtosis:                15.283     Cond. No.:                493.
=====
```

Warnings:

```
[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.
```

Thank you for your attention!

Some links:

- <https://bolighed.dk/>
- <https://da-dk.facebook.com/bolighed/>
- <https://twitter.com/bolighed>

